

### Task 3: Customer Segmentation / Clustering

Plan for Task 3:

Prepare Data:

Combine profile information (from Customers.csv) and transaction data (from Transactions.csv).

Engineer meaningful features for clustering.

1.Select a Clustering Algorithm:

Use K-Means or any other clustering algorithm (like DBSCAN or Agglomerative Clustering).

Define a range for the number of clusters (2 to 10) and evaluate the optimal one.

2.Evaluate Clustering:

Calculate Davies-Bouldin (DB) Index.

Optionally calculate metrics like silhouette score or inertia.

3.Visualize Clusters:

Use scatter plots, pair plots, or cluster centers to visualize.

4.Deliverables:

Number of clusters formed.

DB Index value. Clustering visualization(this one comes under other matrices).

Step 1: Prepare Data for Clustering

```
In [13]: from sklearn.preprocessing import StandardScaler
import pandas as pd
```

```
In [14]: # Load datasets
customers = pd.read_csv("Customers.csv")
transactions = pd.read_csv("Transactions.csv")
```

```
In [15]: # Merge datasets
merged_data = transactions.merge(customers, on="CustomerID", how="left")
```

```
In [16]: # Feature engineering for clustering
customer_clustering_features = merged_data.groupby('CustomerID').agg(
    total_spending=('TotalValue', 'sum'),
    total_transactions=('TransactionID', 'count'),
    avg_transaction_value=('TotalValue', 'mean')
).reset_index()

print(customer_clustering_features)
```

	CustomerID	total_spending	total_transactions	avg_transaction_value
0	C0001	3354.52	5	670.904000
1	C0002	1862.74	4	465.685000
2	C0003	2725.38	4	681.345000
3	C0004	5354.88	8	669.360000
4	C0005	2034.24	3	678.080000
..	...	...	...	...
194	C0196	4982.88	4	1245.720000
195	C0197	1928.65	3	642.883333
196	C0198	931.83	2	465.915000
197	C0199	1979.28	4	494.820000
198	C0200	4758.60	5	951.720000

[199 rows x 4 columns]

```
In [17]: # Add profile information (e.g., Region)
customer_clustering_features = customer_clustering_features.merge(
    customers[['CustomerID', 'Region']], on='CustomerID', how='left'
)
print(customers[['CustomerID', 'Region']].head())
```

	CustomerID	Region
0	C0001	South America
1	C0002	Asia
2	C0003	South America
3	C0004	South America
4	C0005	Asia

```
In [18]: # Check the columns of customer_clustering_features
print(customer_clustering_features.columns)
```

```
# Verify if 'Region' exists
if 'Region' not in customer_clustering_features.columns:
    print("Region column is missing from the DataFrame.")
```

```
Index(['CustomerID', 'total_spending', 'total_transactions',
      'avg_transaction_value', 'Region'],
      dtype='object')
```

```
In [19]: # One-hot encode 'Region'
customer_clustering_features = pd.get_dummies(customer_clustering_features, columns=['Region'])
# successful
```

```
In [20]: # Standardize features

# Step 1: Drop unnecessary or duplicate columns
columns_to_drop = ['Region_x', 'Region_y']
customer_clustering_features = customer_clustering_features.drop(columns=columns_to_drop, errors='ignore')
```

```
In [21]: # Step 2: Remove duplicate one-hot-encoded columns
customer_clustering_features = customer_clustering_features.loc[:, ~customer_clustering_features.columns.duplicated()]
```

```
In [22]: # Step 3: Ensure all features are numeric
clustering_features = customer_clustering_features.drop(columns=['CustomerID'])
clustering_features = pd.get_dummies(clustering_features, drop_first=True)
```

```
In [23]: # Step 4: Standardize the numerical features
scaler = StandardScaler()
clustering_features_scaled = scaler.fit_transform(clustering_features)
```

```
In [24]: print(clustering_features.dtypes)
```

```
total_spending      float64
total_transactions   int64
avg_transaction_value float64
Region_Asia          bool
Region_Europe        bool
Region_North America bool
Region_South America bool
dtype: object
```

```
In [25]: # Step 5: Verify the cleaned and scaled features
print(clustering_features_scaled[:5])
```

```
[[-0.06170143 -0.01145819 -0.07026341 -0.53279543 -0.57928445 -0.54831888
  1.54041597]
 [-0.87774353 -0.46749414 -0.93493297  1.87689298 -0.57928445 -0.54831888
 -0.6491753 ]
 [-0.40585722 -0.46749414 -0.02627131 -0.53279543 -0.57928445 -0.54831888
  1.54041597]
 [ 1.03254704  1.35664965 -0.0767689  -0.53279543 -0.57928445 -0.54831888
  1.54041597]
 [-0.78392861 -0.92353008 -0.04002806  1.87689298 -0.57928445 -0.54831888
 -0.6491753  ]]
```

We completed the code to prepare the data set successfully

### **\*Lets start clustering\***

**\*We can use different algorithms for clustering (e.g., K-Means, DBSCAN, or Agglomerative Clustering)\*** For this task, K-Means Clustering is a great choice because:

Simplicity: It's straightforward and easy to implement for customer segmentation.

Interpretability: K-Means provides well-defined clusters, and centroids can give insights about cluster characteristics.

DB Index Evaluation: K-Means is ideal for optimizing the Davies-Bouldin Index (DB Index) since it focuses on minimizing within-cluster variance.

Task:

Run K-Means Clustering with a range of cluster counts (from 2 to 10).

Calculate Clustering Metrics:

Davies-Bouldin Index (DB Index).

Inertia (Within-cluster Sum of Squares).

Visualize Clusters using scatter plots and possibly a heatmap for better understanding.

```
In [26]: from sklearn.cluster import KMeans
from sklearn.metrics import davies_bouldin_score
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
```

```
In [27]: import warnings

warnings.filterwarnings("ignore", category=UserWarning)
```

```
In [28]: # Range for number of clusters
cluster_range = range(2, 11)
```

```
In [29]: # To store metrics for different cluster sizes/ loop for calculating metrics
db_indices = []
inertia_scores = []
```

```
In [30]: # Run K-Means for each cluster size and calculate metrics
for n_clusters in cluster_range:
    kmeans = KMeans(n_clusters=n_clusters, random_state=42, n_init=10)
    kmeans.fit(clustering_features_scaled)

    # Cluster Labels
    labels = kmeans.labels_
```

```
# Calculate DB Index and Inertia
db_index = davies_bouldin_score(clustering_features_scaled, labels)
inertia = kmeans.inertia_

# Append one value for each cluster size
db_indices.append(db_index)
inertia_scores.append(inertia)
```

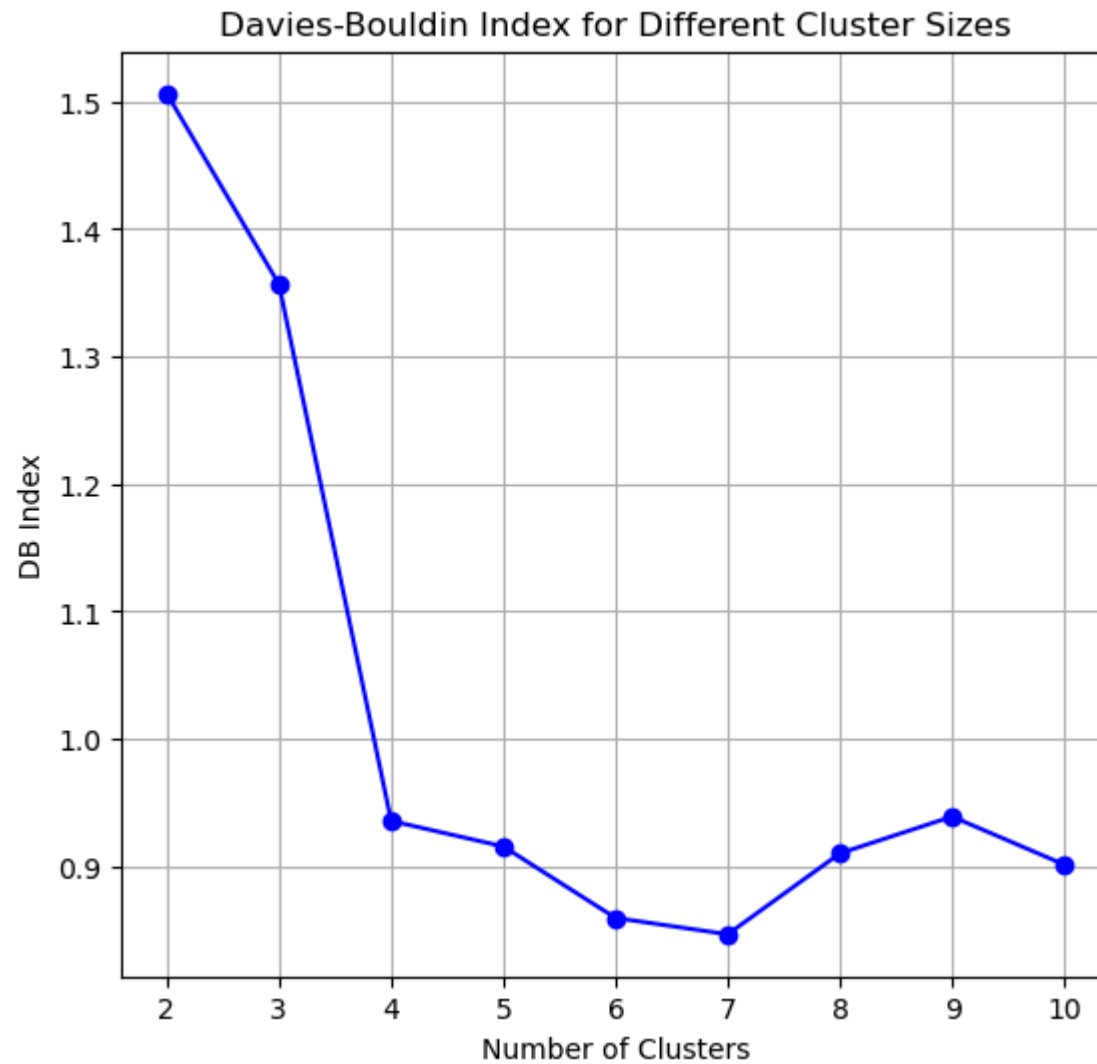
```
In [31]: # Confirm lengths of db_indices and inertia_scores match cluster_range
print("Length of cluster_range:", len(cluster_range))
print("Length of db_indices:", len(db_indices))
print("Length of inertia_scores:", len(inertia_scores))
```

```
Length of cluster_range: 9
Length of db_indices: 9
Length of inertia_scores: 9
```


```
In [32]: # Plot metrics to choose the best number of clusters
plt.figure(figsize=(14, 6))
```

```
Out[32]: <Figure size 1400x600 with 0 Axes>
<Figure size 1400x600 with 0 Axes>
```

```
In [33]: # Plot Davies-Bouldin Index
plt.figure(figsize=(14, 6))
plt.subplot(1, 2, 1)
plt.plot(cluster_range, db_indices, marker='o', linestyle='--', color='blue')
plt.title("Davies-Bouldin Index for Different Cluster Sizes")
plt.xlabel("Number of Clusters")
plt.ylabel("DB Index")
plt.grid(True)
```



```
In [34]: # Plot Inertia
plt.subplot(1, 2, 2)
plt.plot(cluster_range, inertia_scores, marker='o', linestyle='-', color='green')
plt.title("Inertia for Different Cluster Sizes")
plt.xlabel("Number of Clusters")
plt.ylabel("Inertia")
plt.grid(True)
```



```
In [35]: plt.tight_layout()  
plt.show()
```

<Figure size 640x480 with 0 Axes>

```
In [37]: # Choose the optimal number of clusters based on the above plots  
optimal_clusters = int(input("Enter the optimal number of clusters based on the plots: "))
```

Enter the optimal number of clusters based on the plots: 4

```
In [38]: # Run K-Means with the optimal number of clusters  
final_kmeans = KMeans(n_clusters=optimal_clusters, random_state=42, n_init=10)  
final_kmeans.fit(clustering_features_scaled)
```



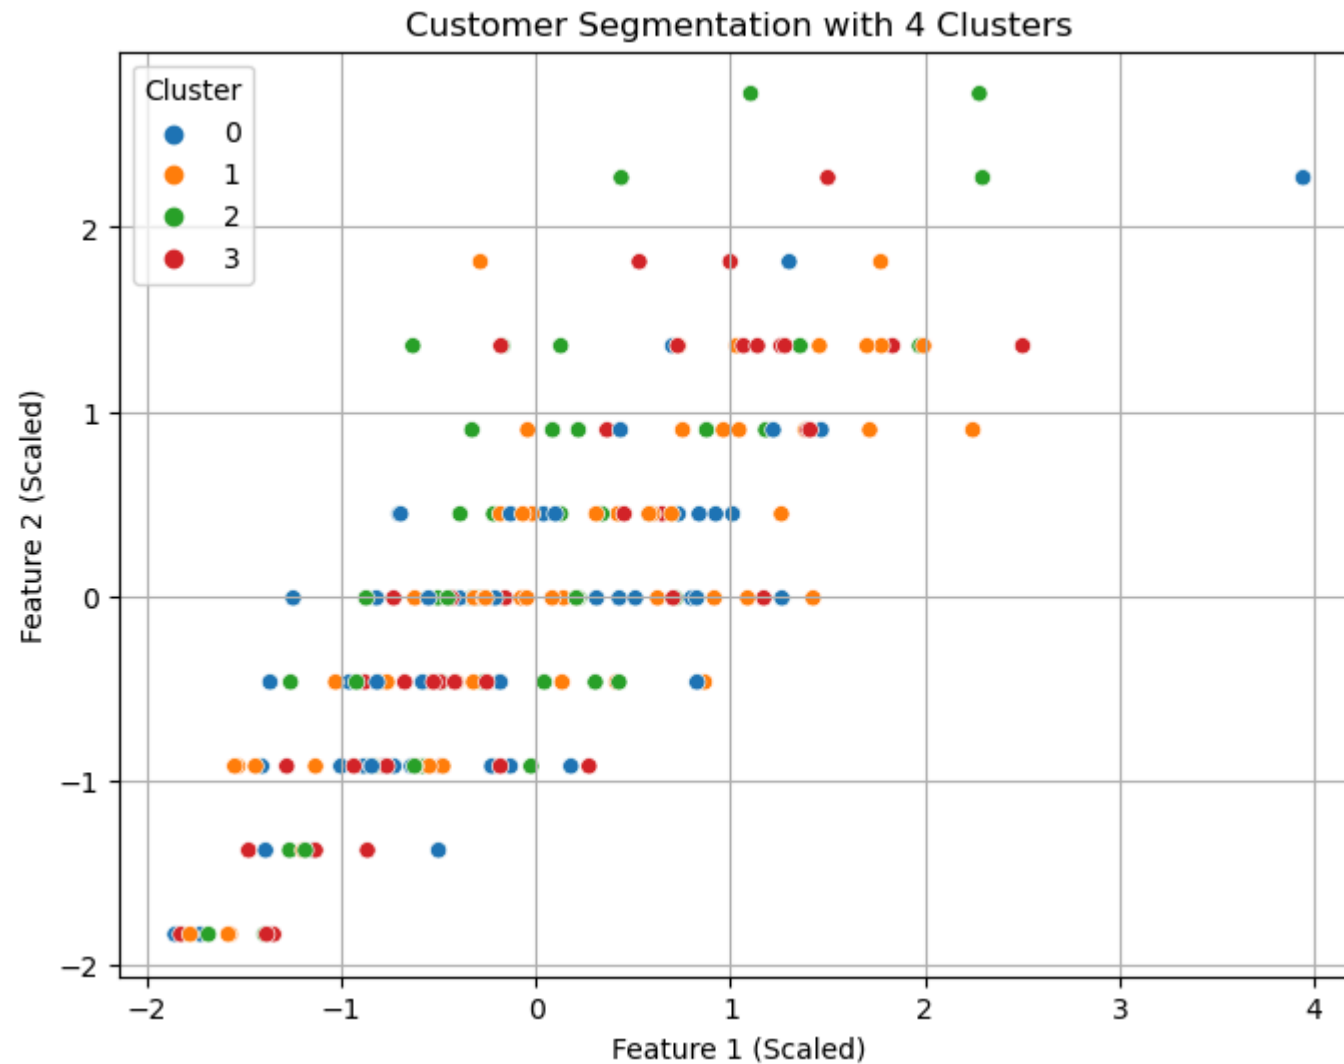
Out[38]:

▼ KMeans

```
KMeans(n_clusters=4, n_init=10, random_state=42)
```

```
In [39]: # Add the cluster labels to the original dataset
customer_clustering_features['Cluster'] = final_kmeans.labels_
```

```
In [40]: # Visualize final clusters (using the first two features for simplicity)
plt.figure(figsize=(8, 6))
sns.scatterplot(
    x=clustering_features_scaled[:, 0],
    y=clustering_features_scaled[:, 1],
    hue=final_kmeans.labels_,
    palette="tab10",
    legend="full"
)
plt.title(f"Customer Segmentation with {optimal_clusters} Clusters")
plt.xlabel("Feature 1 (Scaled)")
plt.ylabel("Feature 2 (Scaled)")
plt.legend(title="Cluster")
plt.grid(True)
plt.show()
```



```
In [41]: # Print final DB Index for the chosen clusters
final_db_index = davies_bouldin_score(clustering_features_scaled, final_kmeans.labels_)
print(f"Final Davies-Bouldin Index for {optimal_clusters} clusters: {final_db_index}")
```

Final Davies-Bouldin Index for 4 clusters: 0.9355298648489481

By this Task 3: Customer Segmentation / Clustering completed

- Clustering logic and metrics.
- Visual representation of clusters.